

Intro to DB

# CHAPTER 2

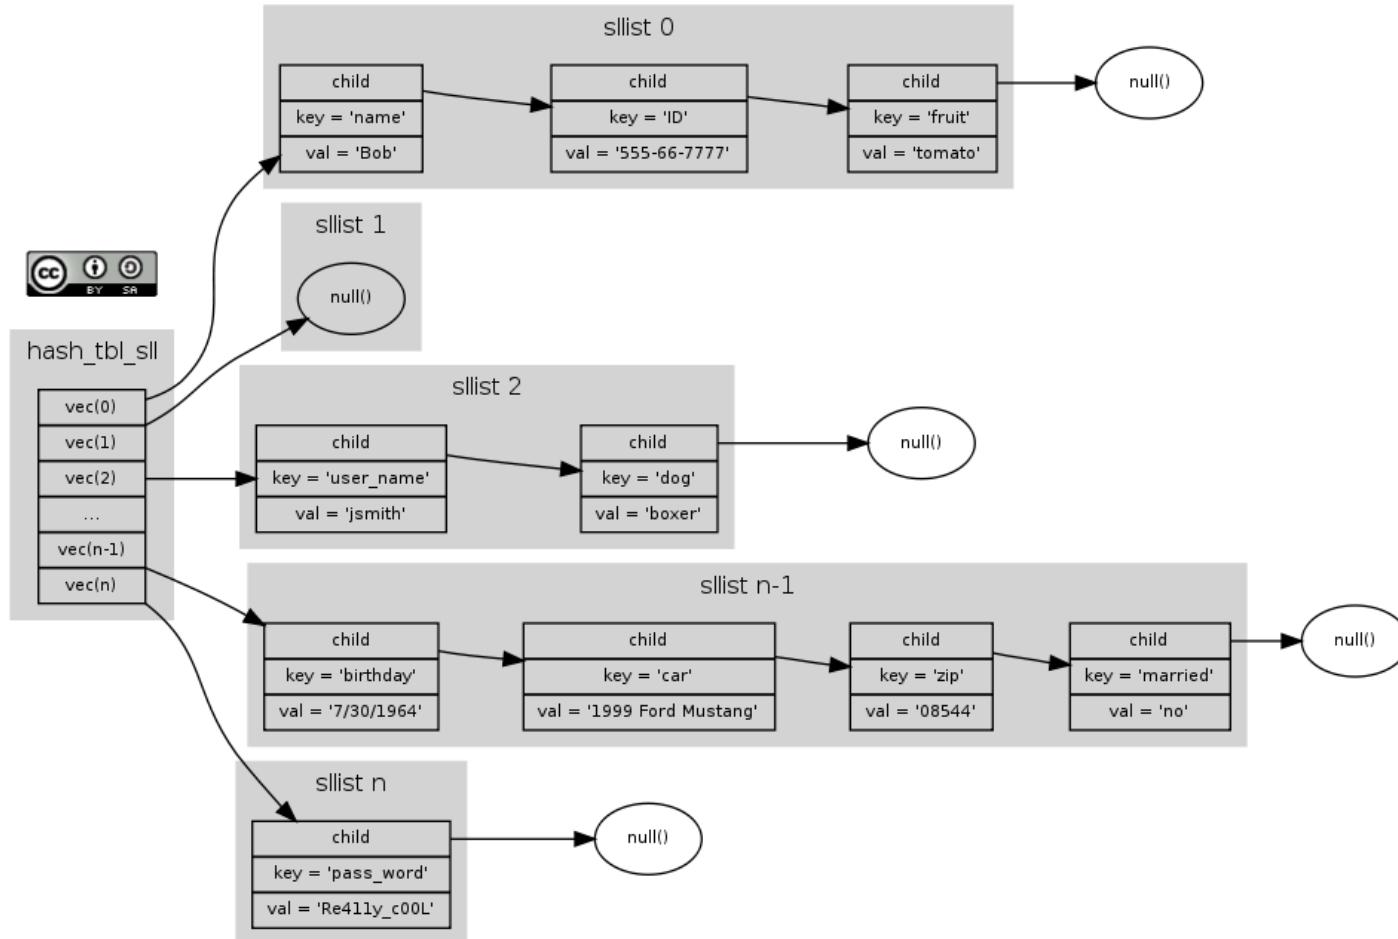
# RELATIONAL MODEL

# Chapter 2: Relational Model

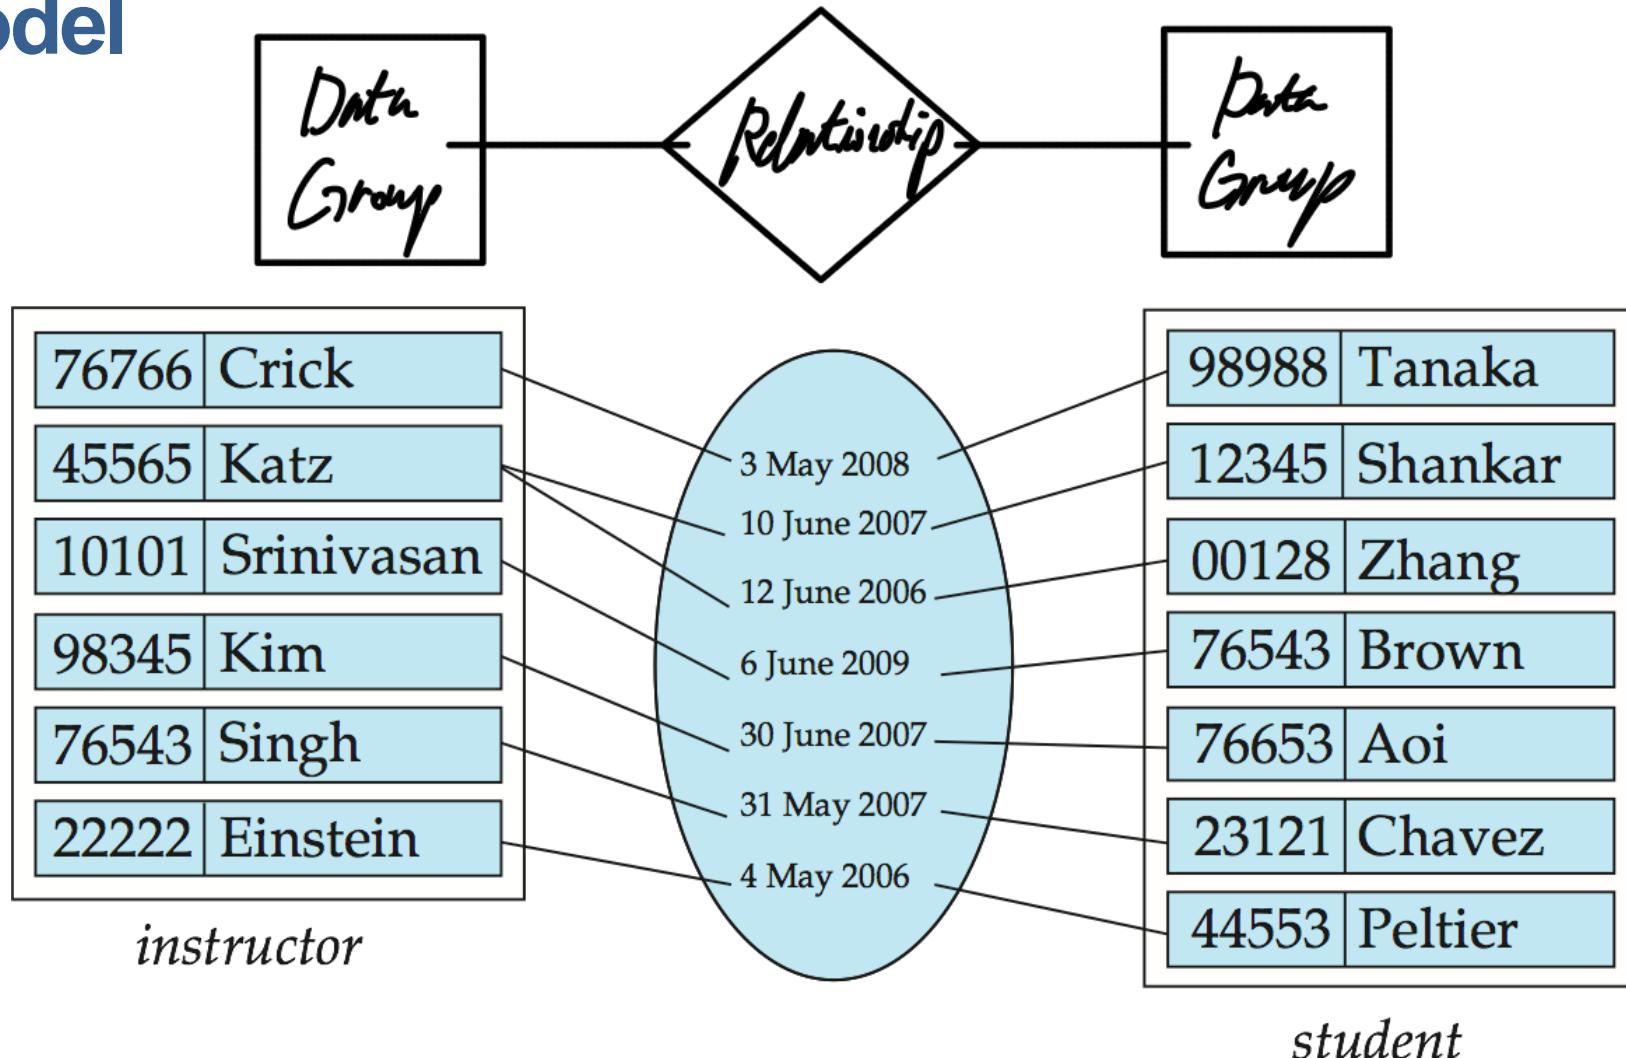
- Structure of Relational Databases
- Database Schema
- Keys
- Schema Diagrams
- Relational Query Languages
- Relational Algebra

# Data Model

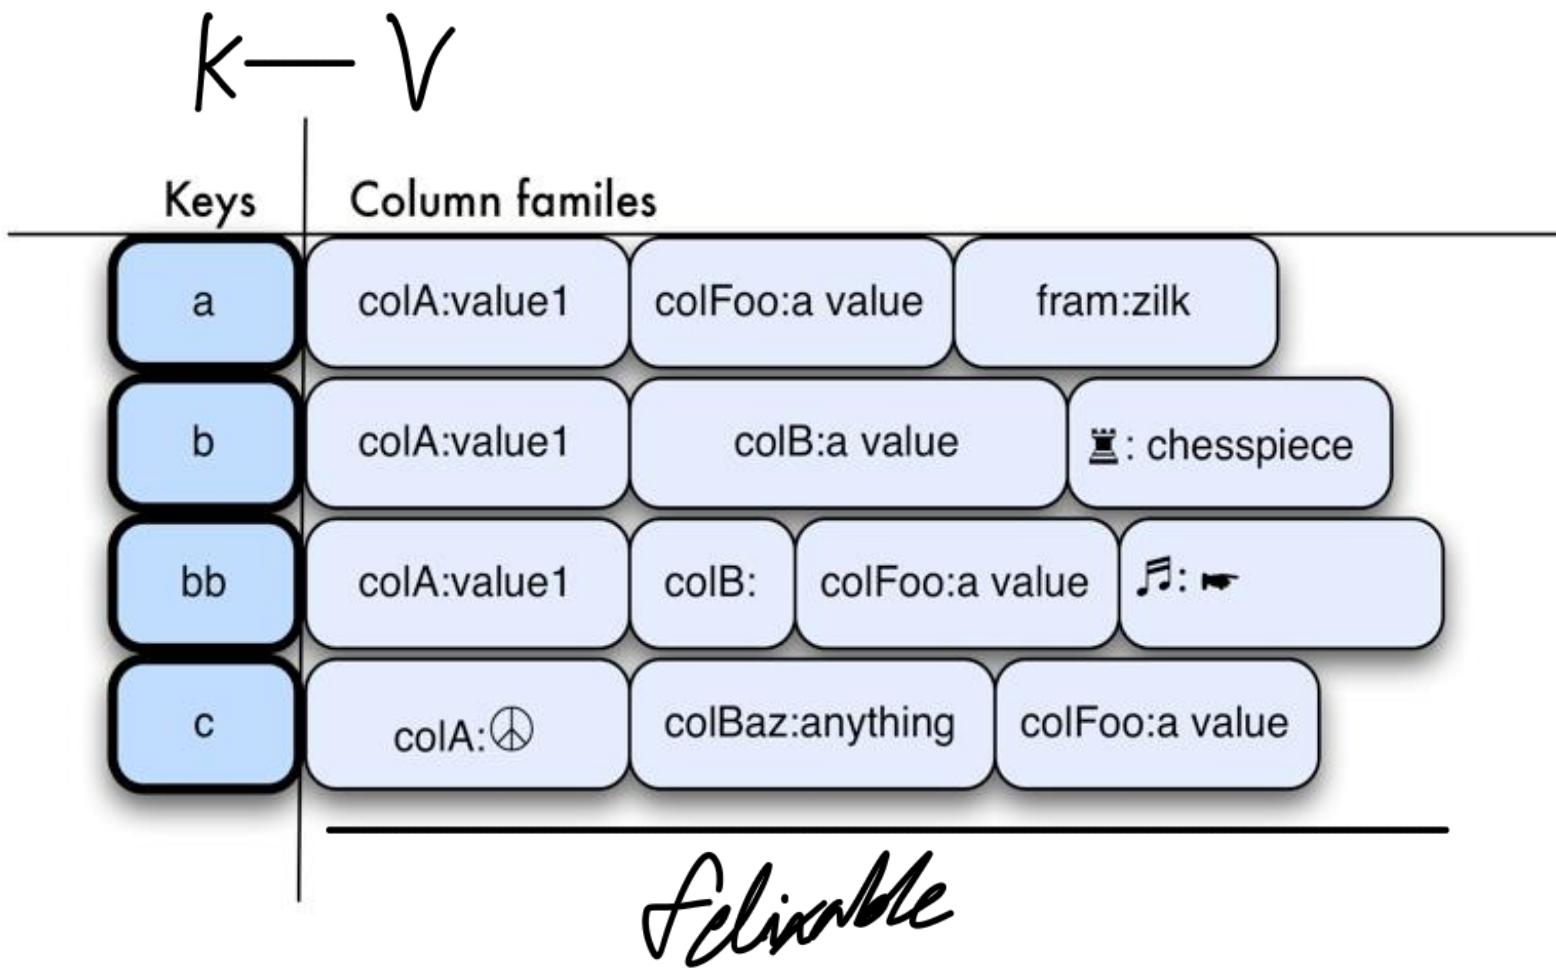
- The framework/formalism for representing data and their relationships



# Data Model



# Data Model



# Data Model

Column

The diagram shows a Microsoft Excel window titled "Sheet1 - Microsoft Excel". The window displays a table with 26 rows and 11 columns. The columns are labeled A through J, and the rows are numbered 1 through 26. The first row contains column headers: EMP\_ID, JOB\_ID, DPT\_MANAGER, EMP\_DEPT, LOCATION\_ID, COUNTRY\_ID, EMP\_NAME, EMP\_SALARY, COMMISSION\_PCT, and DPT\_NAME. The data rows follow, showing various employee details such as name, salary, and department. Handwritten labels are present: "Row" with an arrow pointing down the first column, "Column" with an arrow pointing right across the first row, and "Cell" with an arrow pointing to the cell containing "EDWARD JOSEPH" in row 18, column B.

EMP_ID	JOB_ID	DPT_MANAGER	EMP_DEPT	LOCATION_ID	COUNTRY_ID	EMP_NAME	EMP_SALARY	COMMISSION_PCT	DPT_NAME
1	77514063	AC_ACCOUNT	48748339	LI	6	KATHY ABELLERA	650166.01	0.44	Library Interloan (ext 6514) (see Library Information Services Team)
2	21328077	AC_ACCOUNT	48748339	LI	6	JEREMIAH DINETZ	186958.24	0.26	Library Interloan (ext 6514) (see Library Information Services Team)
3	63114872	AC_ACCOUNT	48748339	LI	6	WENDALEE BURNETT	683337.14	0.47	Library Interloan (ext 6514) (see Library Information Services Team)
4	22447298	AC_ACCOUNT	48748339	LI	6	MISSY KOWAL	413672.3	0.06	Library Interloan (ext 6514) (see Library Information Services Team)
5	42823120	AC_ACCOUNT	48748339	LI	6	JEHU POWELL	878187.64	0.44	Library Interloan (ext 6514) (see Library Information Services Team)
6	73022131	AC_ACCOUNT	48748339	LI	6	BAZE WALTON	231658.56	0.16	Library Interloan (ext 6514) (see Library Information Services Team)
7	30206729	AC_ACCOUNT	48748339	LI	6	DEE TOOF	593010.32	0.6	Library Interloan (ext 6514) (see Library Information Services Team)
8	68388930	AC_ACCOUNT	48748339	LI	6	MURRAY WALEN	274527.68	0.56	Library Interloan (ext 6514) (see Library Information Services Team)
9	52994681	AC_ACCOUNT	48748339	LI	6	TODD CORNELL	366530.46	0.59	Library Interloan (ext 6514) (see Library Information Services Team)
10	85531859	AC_ACCOUNT	48748339	LI	6	CRAIG HARRINGTON	207694.73	0.14	Library Interloan (ext 6514) (see Library Information Services Team)
11	99625806	AC_ACCOUNT	48748339	LI	6	DALIA DEROSIER	788803.05	0.44	Library Interloan (ext 6514) (see Library Information Services Team)
12	85952542	AC_ACCOUNT	48748339	LI	6	JERAMY CAMPISE	92348.6	0.48	Library Interloan (ext 6514) (see Library Information Services Team)
13	90514846	AC_ACCOUNT	48748339	LI	6	ERV BRODSHO	674926	0.46	Library Interloan (ext 6514) (see Library Information Services Team)
14	10293371	AC_ACCOUNT	48748339	LI	6	RODGER WAKEFIELD	521169.69	0.6	Library Interloan (ext 6514) (see Library Information Services Team)
15	93331872	AC_ACCOUNT	48748339	LI	6	GEODIE ROEHL	607693.22	0.7	Library Interloan (ext 6514) (see Library Information Services Team)
16	38996915	AC_ACCOUNT	48748339	LI	6	BRUCE BULAGA	346979.22	0.68	Library Interloan (ext 6514) (see Library Information Services Team)
17	95488542	AC_ACCOUNT	48748339	LI	6	EDWARD JOSEPH	405804.03	0.39	Library Interloan (ext 6514) (see Library Information Services Team)
18	43800025	AC_ACCOUNT	48748339	LI	6	BRIAN FURNO	652747.04	0.3	Library Interloan (ext 6514) (see Library Information Services Team)
19	22735685	AC_ACCOUNT	48748339	LI	6	VERMELL DRESSER	426205.82	0.71	Library Interloan (ext 6514) (see Library Information Services Team)
20	11249307	AC_ACCOUNT	48748339	LI	6	EILEEN GILBERTSON	600265.45	0.27	Library Interloan (ext 6514) (see Library Information Services Team)
21	30157047	AC_ACCOUNT	48748339	LI	6	ANGELIA HOWETH	94186.38	0.33	Library Interloan (ext 6514) (see Library Information Services Team)
22	62528105	AC_ACCOUNT	48748339	LI	6	KASEY STRAMANDINOLI	718647.7	0.57	Library Interloan (ext 6514) (see Library Information Services Team)
23	64880465	AC_ACCOUNT	48748339	LI	6	GRAIG DOWNER	618324.31	0.5	Library Interloan (ext 6514) (see Library Information Services Team)
24	68221632	AC_ACCOUNT	48748339	LI	6	OCTAVIA LAMOREAUX	372831.78	0.37	Library Interloan (ext 6514) (see Library Information Services Team)
25	87701077	AC_ACCOUNT	48748339	LI	6	AUTUMN BOLTON	531229.73	0.63	Library Interloan (ext 6514) (see Library Information Services Team)
26	32060000	AC_ACCOUNT	48748339	LI	6	BOBNET BOUCHER	530760.91	0.67	Library Interloan (ext 6514) (see Library Information Services Team)

# Data Model - Relational

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	
10101	Srinivasan	Comp. Sci.	65000	
12121	Wu	Finance	90000	
15151	Mozart	Music	40000	
22222	Einstein	Physics	95000	
32343	El Said	History	60000	
33456	Gold	Physics	87000	
45565	Katz	Comp. Sci.	75000	
58583	Califieri	History	62000	
76543	Singh	Finance	80000	
76766	Crick	Biology	72000	
83821	Brandt	Comp. Sci.	92000	
98345	Kim	Elec. Eng.	80000	

Diagram illustrating the relational data model:

- attributes (or columns):** Indicated by arrows pointing to the column headers *ID*, *name*, *dept\_name*, and *salary*.
- tuples (or rows):** Indicated by arrows pointing to the data rows in the table.
- passive data (not what you just do):** A handwritten note with an arrow pointing to the last row of the table.

# Relation (in Set Theory)

- **Binary relation**  $R$  on a set  $A$   
is a collection of ordered pairs of elements of  $A$

$$A = \{1, 2, 3\}$$

- $R \subseteq A \times A$

$$< = \{\langle 1, 2 \rangle, \langle 1, 3 \rangle, \langle 2, 3 \rangle\}$$

$$\leq = \{\langle 1, 1 \rangle, \langle 1, 2 \rangle, \langle 1, 3 \rangle, \langle 2, 2 \rangle, \langle 2, 3 \rangle, \langle 3, 3 \rangle\}$$

- (*Bipartite*) Relation  $R_2$  between elements of set  $B$  and set  $C$

ex)  $B = \text{Set of Country} / C = \text{Set of City}$

$$\text{Capital City} = \{\langle B_i, C_j \rangle \dots\}$$

# Basic Structure of a Relation

- Formally,

given sets  $D_1, D_2, \dots, D_n$ , a *relation*  $r$  is a subset of  $D_1 \times D_2 \times \dots \times D_n$

$$r \subseteq D_1 \times \dots \times D_n \quad (n: \text{degree of } r)$$

or

$$r = \{ \langle d_1, \dots, d_n \rangle \mid d_1 \in D_1, \dots, d_n \in D_n \} \quad (\text{set of tuples})$$

- Example:  
 $\text{customer-name} = \{\text{Jones, Smith, Curry, Lindsay}\}$   
 $\text{customer-street} = \{\text{Main, North, Park}\}$   
 $\text{customer-city} = \{\text{Harrison, Rye, Pittsfield}\}$

Then,

$$r = \{ (\text{Jones, Main, Harrison}), (\text{Smith, North, Rye}), (\text{Curry, North, Rye}), (\text{Lindsay, Park, Pittsfield}) \}$$

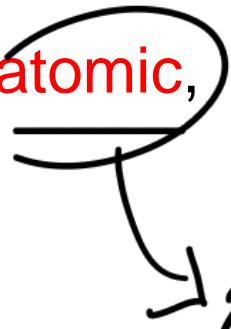
is a relation over  $\text{customer-name} \times \text{customer-street} \times \text{customer-city}$

*actual data of customer's address*

*Attributes*

# Attribute Types

- Each attribute of a relation has a name
- **Domain** of an attribute: the set of allowed values allowed for each attribute
- Attribute values are (normally) required to be **atomic**,  
that is, indivisible
  - E.g. multivalued attribute values are not atomic
  - E.g. composite attribute values are not atomic
- The special value **null** is a member of every domain
  - Represents unknown or missing value



*minimal unit of  
each data*

# Relation Schema

*Set*

- $A_1, A_2, \dots, A_n$  are *attributes*
- $R = (A_1, A_2, \dots, A_n)$  is a relation schema

e.g.

Customer-schema = (*customer-name*, *customer-street*, *customer-city*)

*schema*

- $r(R)$  is a relation (variable) on the *relation schema*  $R$

e.g.

*customer*(Customer-schema)

*relation Schema*(*type*)

ex)  $R$ : *type*

$r$ : *variable name*

# Relation Instance

- The current values of a relation
  - Specified by a table
- An element  $t$  of  $r$  is a *tuple*,
  - represented by a *row* in a table

*Instance*

The diagram shows a table with three columns labeled *customer-name*, *customer-street*, and *customer-city*. The table has four rows of data. A handwritten-style arrow labeled "Instance" points from the top-left of the table area towards the top right. Another arrow labeled "attributes" points from the column headers to the top edge of the table. A third arrow labeled "tuples" points from the bottom center of the table to the right edge. The data in the table is as follows:

customer-name	customer-street	customer-city
Jones	Main	Harrison
Smith	North	Rye
Curry	North	Rye
Lindsay	Park	Pittsfield

*customer*

# Relations are Unordered

- Order of tuples is irrelevant  
(tuples may be stored in an arbitrary order)
- E.g. *instructor* relation with unordered tuples

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

# Relational Database : Set of Multiple Table

- A database consists of multiple relations
- Information about an enterprise is broken up into parts, with each relation storing one part of the information
  - E.g.: *instructor*: information about teachers in the university
  - student*: information about students
  - advisor*: information about which instructor advises which students
- Storing all information as a single relation is not a good idea
  - univ (instructor -ID, name, dept\_name, salary, student\_Id, ...)*
- Database design (Chapter 6 & 7) deals with how to decide on the relational schemas

# A relational database for a university

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	K. S.	Computer Sci.	75000

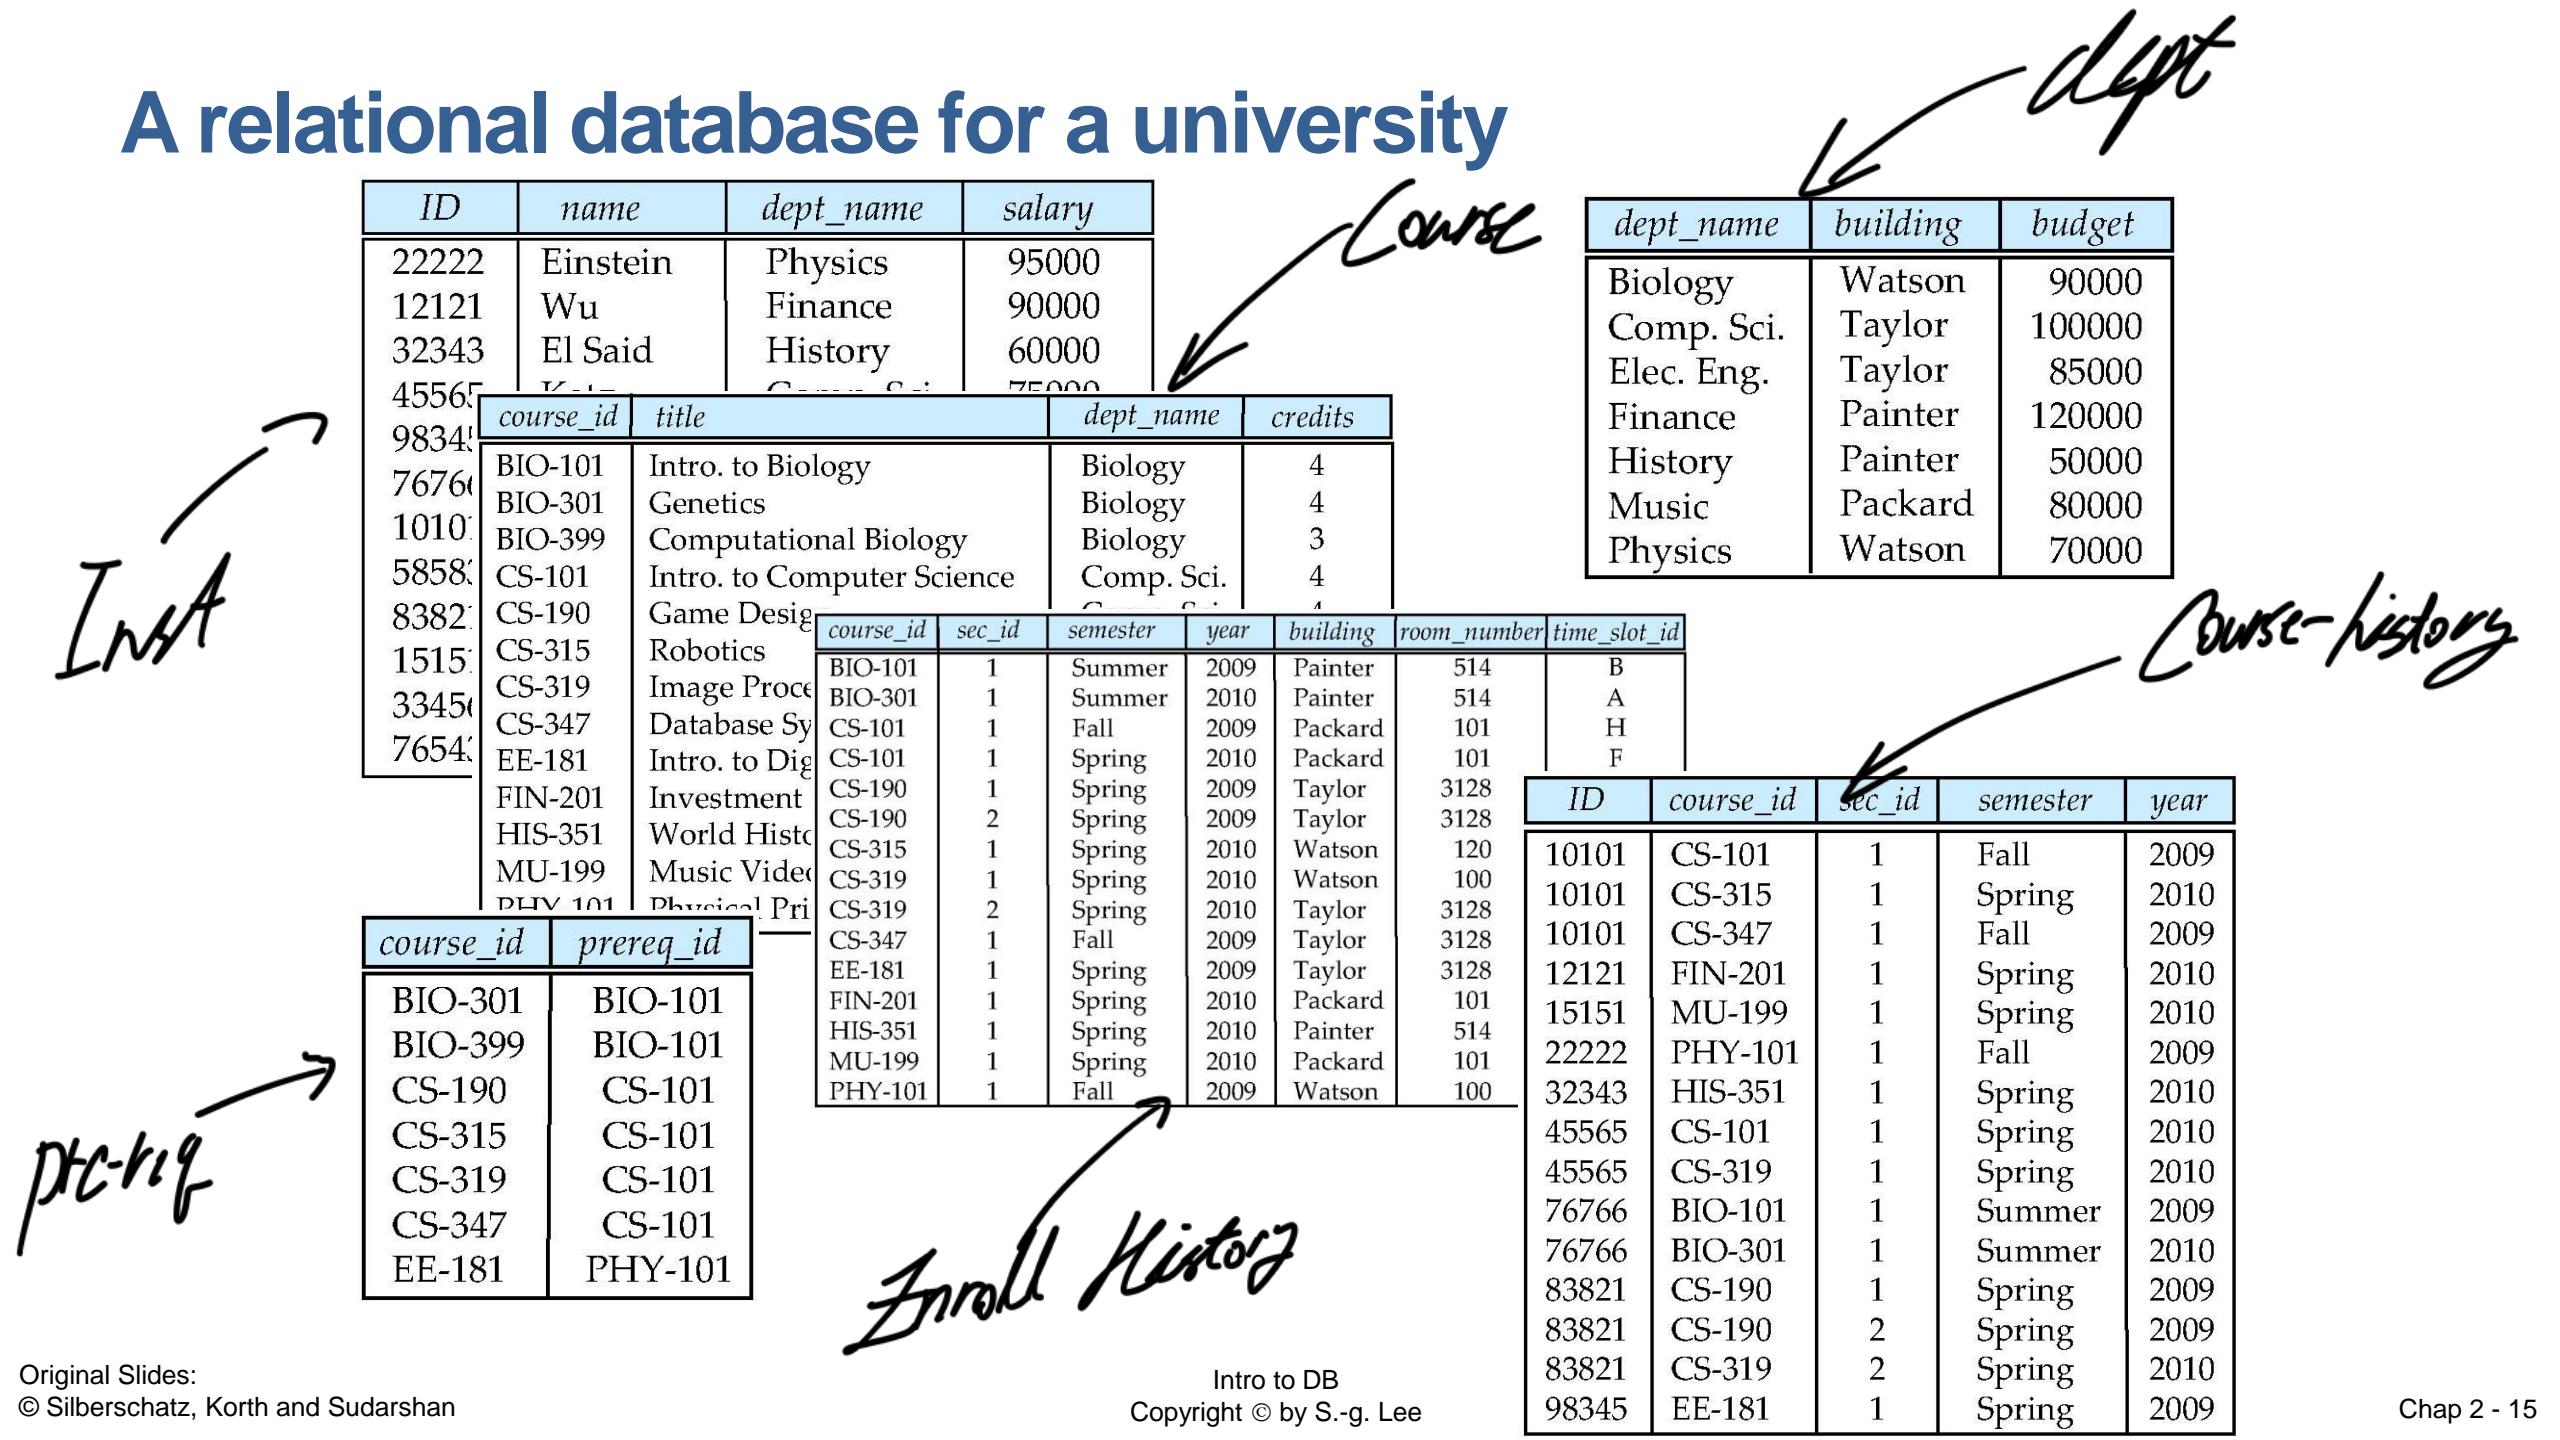
course_id	prereq_id
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

course_id	title	dept_name	credits
98345	BIO-101	Intro. to Biology	4
76766	BIO-301	Genetics	4
10101	BIO-399	Computational Biology	3
58585	CS-101	Intro. to Computer Science	4
83822	CS-190	Game Design	4
15151	CS-315	Robotics	4
33456	CS-319	Image Proce	4
76545	CS-347	Database Sy	4
EE-181	EE-181	Intro. to Dig	4
FIN-201	FIN-201	Investment	4
HIS-351	HIS-351	World Histo	4
MU-199	MU-199	Music Video	4
PHY-101	PHY-101	Physical Pri	4

course_id	sec_id	semester	year	building	room_number	time_slot_id
BIO-101	1	Summer	2009	Painter	514	B
BIO-301	1	Summer	2010	Painter	514	A
CS-101	1	Fall	2009	Packard	101	H
CS-101	1	Spring	2010	Packard	101	F
CS-190	1	Spring	2009	Taylor	3128	
CS-190	2	Spring	2009	Taylor	3128	
CS-315	1	Spring	2010	Watson	120	
CS-319	1	Spring	2010	Watson	100	
CS-319	2	Spring	2010	Taylor	3128	
CS-347	1	Fall	2009	Taylor	3128	
EE-181	1	Spring	2009	Taylor	3128	
FIN-201	1	Spring	2010	Packard	101	
HIS-351	1	Spring	2010	Painter	514	
MU-199	1	Spring	2010	Packard	101	
PHY-101	1	Fall	2009	Watson	100	

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009
32343	HIS-351	1	Spring	2010
45565	CS-101	1	Spring	2010
45565	CS-319	1	Spring	2010
76766	BIO-101	1	Summer	2009
76766	BIO-301	1	Summer	2010
83821	CS-190	1	Spring	2009
83821	CS-190	2	Spring	2009
83821	CS-319	2	Spring	2010
98345	EE-181	1	Spring	2009



## Keys

*K is single or combination of Attributes*

- Let  $K \subseteq R$  *unique value for each instance*
- $K$  is a superkey of  $R$   
if values for  $K$  are sufficient to identify a unique tuple of each possible relation  $r(R)$ .
- By “possible  $r$ ” we mean a relation  $r$  that could exist in the enterprise we are modeling.
  - Example:  $\{\text{customer-name}, \text{customer-street}\}$  and  $\{\text{customer-name}\}$  are both superkeys of *Customer*, if no two customers can possibly have the same name.
- $K$  is a candidate key if  $K$  is minimal
  - Example:  $\{\text{customer-name}\}$  is a candidate key for *Customer*, since it is a superkey, and no subset of it is a superkey.

# Keys

- **Primary key:** one of the candidate keys is selected to be the primary key of the table

- which one? : *representative Candidate key of Data*

*primary key*

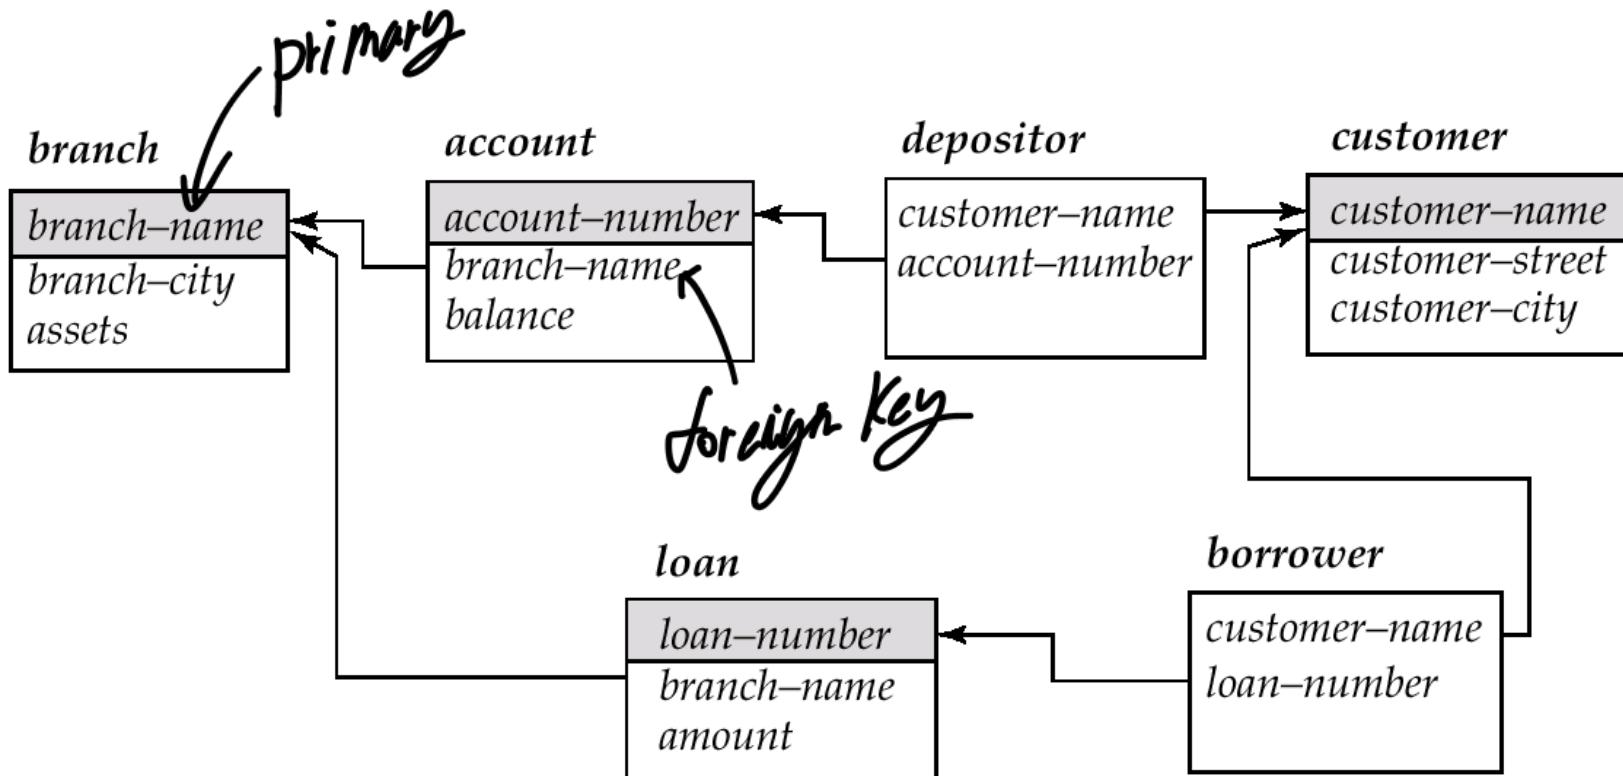
- **Foreign key:** Value in one relation must appear in another

- **Referencing** relation
- **Referenced** relation

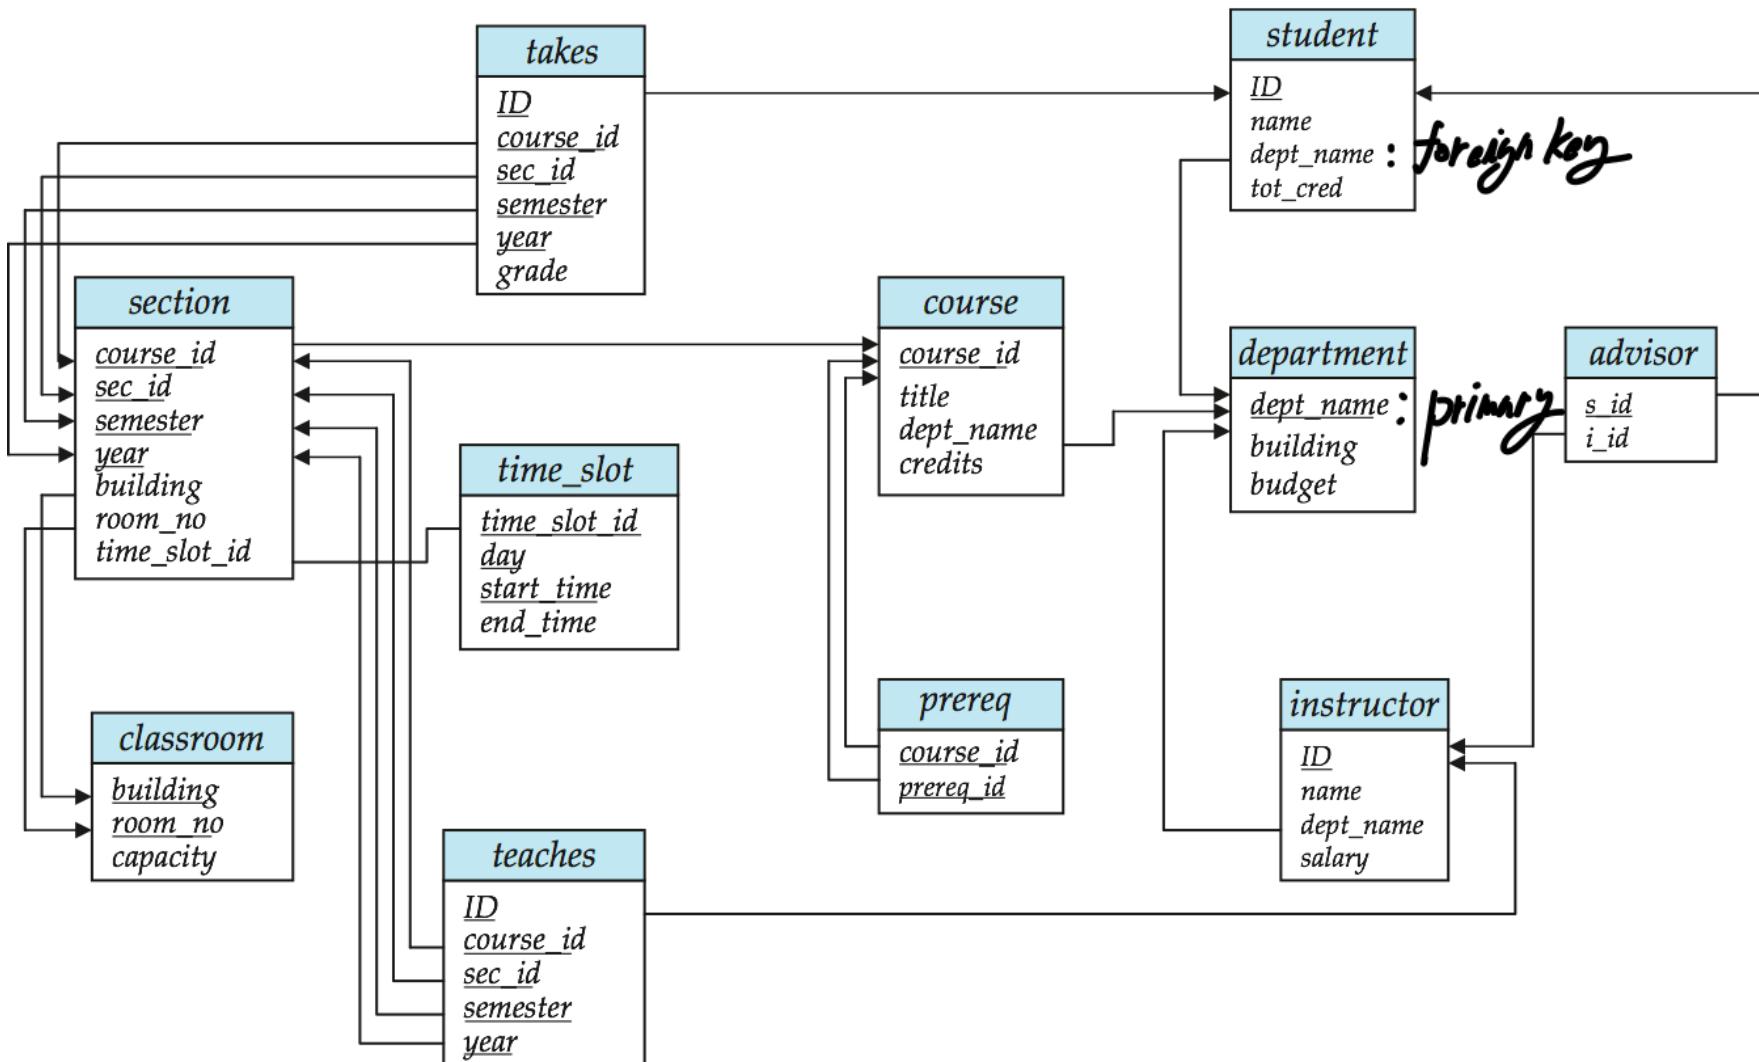
: *kind of Pointer*

course_id	prereq_id
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

# Schema Diagram for a Banking Enterprise



# Schema Diagram for the University



# Query Languages

- Language in which user requests information from the database.
  - procedural : *Sequential*
    - specify what data are needed
    - and “how” to get those data
  - nonprocedural: *Conditional (SQL)*
    - specify only “what” data are needed
    - declarative
- “Pure” languages:
  - Relational Algebra
  - Tuple Relational Calculus
  - Domain Relational Calculus
- Pure languages form underlying basis of query languages that people use.

# Relational Algebra : *Procedural*

ERF

- Algebra : operators and operands
  - Relational algebra
    - operands : relations
    - operators : basic operators (+ additional operations)
      - take two or more relations as inputs and give a new relation as a result.
- Procedural language
- Operators
  - select
  - project
  - union
  - set difference
  - Cartesian product
  - rename
  - join
  - division
  - assignment
  - ...

# Examples of Relational Operators

Symbol (Name)	Example of Use
$\sigma$ (Selection)	$\sigma_{\text{salary} \geq 85000}(\text{instructor})$ Return rows of the input relation that satisfy the predicate.
$\Pi$ (Projection)	$\Pi_{ID, \text{salary}}(\text{instructor})$ Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output.
$\bowtie$ (Natural Join)	$\text{instructor} \bowtie \text{department}$ Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.
$\times$ (Cartesian Product)	$\text{instructor} \times \text{department}$ Output all pairs of rows from the two input relations (regardless of whether or not they have the same values on common attributes)
$\cup$ (Union)	$\Pi_{name}(\text{instructor}) \cup \Pi_{name}(\text{student})$ Output the union of tuples from the two input relations.

# Select Operation – selection of tuples

Relation  $r$ :

A	B	C	D
$\alpha$	$\alpha$	1	7
$\alpha$	$\beta$	5	7
$\beta$	$\beta$	12	3
$\beta$	$\beta$	23	10

record

$\sigma_{A=B \wedge D>5}(r)$ :

$(A = B) \wedge (D > 5)$   
and

A	B	C	D
$\alpha$	$\alpha$	1	7
$\beta$	$\beta$	23	10

# Select Operation

- Notation:  $\sigma_p(r)$
- $p$  is called the **selection predicate**
- Defined as:  $\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$ 
  - Where  $p$  is a formula in propositional calculus consisting of **terms** connected by :  $\wedge$  (and),  
 $\vee$  (or),  $\neg$  (not)
  - Each **term** is one of:  
 $\langle\text{attribute}\rangle \ op \ \langle\text{attribute}\rangle$  or  $\langle\text{constant}\rangle$   
where  $op$  is one of:  $=, \neq, >, \geq, <, \leq$
- Example:  $\sigma_{\text{branch-name}=\text{"Perryridge}}(\text{account})$

# Project Operation – selection of columns

- Relation  $r$ :

	A	B	C
$\alpha$	10	1	
$\alpha$	20	1	
$\beta$	30	1	
$\beta$	40	2	

get A, C column

$$\Pi_{A,C}(r)$$

	A	C
$\alpha$	1	
$\alpha$	1	
$\beta$	1	
$\beta$	2	

=

	A	C
$\alpha$	1	
$\beta$	1	

# Project Operation

- Notation:

$$\Pi_{A_1, A_2, \dots, A_k}(r)$$

where  $A_1, A_2$  are attribute names and  $r$  is a relation name.

- The result is defined as the relation of  $k$  columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result
  - since relations are sets
- E.g.: To eliminate the *branch-name* attribute of *account*

$$\Pi_{\underline{\text{account-number, balance}}}(account)$$

*↳ Indicate Attribute to get*

# Union Operation – merging two relations

- Relations  $r, s$ :

A	B
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$

A	B
$\alpha$	2
$\beta$	3

$s$

- $r \cup s$ :

A	B
$\alpha$	1
$\alpha$	2
$\beta$	1
$\beta$	3

# Union Operation

- Notation:  $r \cup s$
- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For  $r \cup s$  to be valid (**Union compatible**)
  - $r, s$  must have the same **arity** (same number of attributes)
  - The attribute domains must be **compatible**  
(e.g., 2nd column of  $r$  deals with the same type of values as does the 2<sup>nd</sup> column of  $s$ )

*same schema*

e.g.: to find all customers with either an account or a loan

$$\Pi_{\underline{\text{customer-name}}}(\text{depositor}) \cup \Pi_{\underline{\text{customer-name}}}(\text{borrower})$$

*But  $\Pi_{\text{customer-name}}(\text{depositor} \cup \text{borrower})$  is invalid*

# Set Difference Operation : 차집합

- Relations  $r, s$ :

A	B
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$

A	B
$\alpha$	2
$\beta$	3

$s$

- $r - s$ :

A	B
$\alpha$	1
$\beta$	1

# Set Difference Operation

- Notation  $r - s$
- Defined as:

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- Set differences must be taken between **compatible** relations.
  - $r$  and  $s$  must have the *same arity*
  - attribute domains of  $r$  and  $s$  must be compatible

# Cartesian-Product Operation – Example

- Relations  $r, s$ :

A	B
$\alpha$	1
$\beta$	2

$r$

C	D	E
$\alpha$	10	a
$\beta$	10	a
$\beta$	20	b
$\gamma$	10	b

$s$

- $r \times s$ :

A	B	C	D	E
$\alpha$	1	$\alpha$	10	a
$\alpha$	1	$\beta$	19	a
$\alpha$	1	$\beta$	20	b
$\alpha$	1	$\gamma$	10	b
$\beta$	2	$\alpha$	10	a
$\beta$	2	$\beta$	10	a
$\beta$	2	$\beta$	20	b
$\beta$	2	$\gamma$	10	b

: Combination set of  $r,s$

# Cartesian-Product Operation

- Notation:  $r \times s$
- Defined as:

$$r \times s = \{ t q \mid t \in r \text{ and } q \in s\}$$

- Assume that attributes of  $r(R)$  and  $s(S)$  are disjoint;  $\rightarrow$  no same attr  
i.e.,  $R \cap S = \emptyset$ .
- If not, renaming of attributes is needed.

# Composition of Operations

- Can build expressions using multiple operations
- Example:  $\sigma_{A=C}(\underline{r \times s})$
- $r \times s$

A	B	C	D	E
$\alpha$	1	$\alpha$	10	a
$\alpha$	1	$\beta$	19	a
$\alpha$	1	$\beta$	20	b
$\alpha$	1	$\gamma$	10	b
$\beta$	2	$\alpha$	10	a
$\beta$	2	$\beta$	10	a
$\beta$	2	$\beta$	20	b
$\beta$	2	$\gamma$	10	b

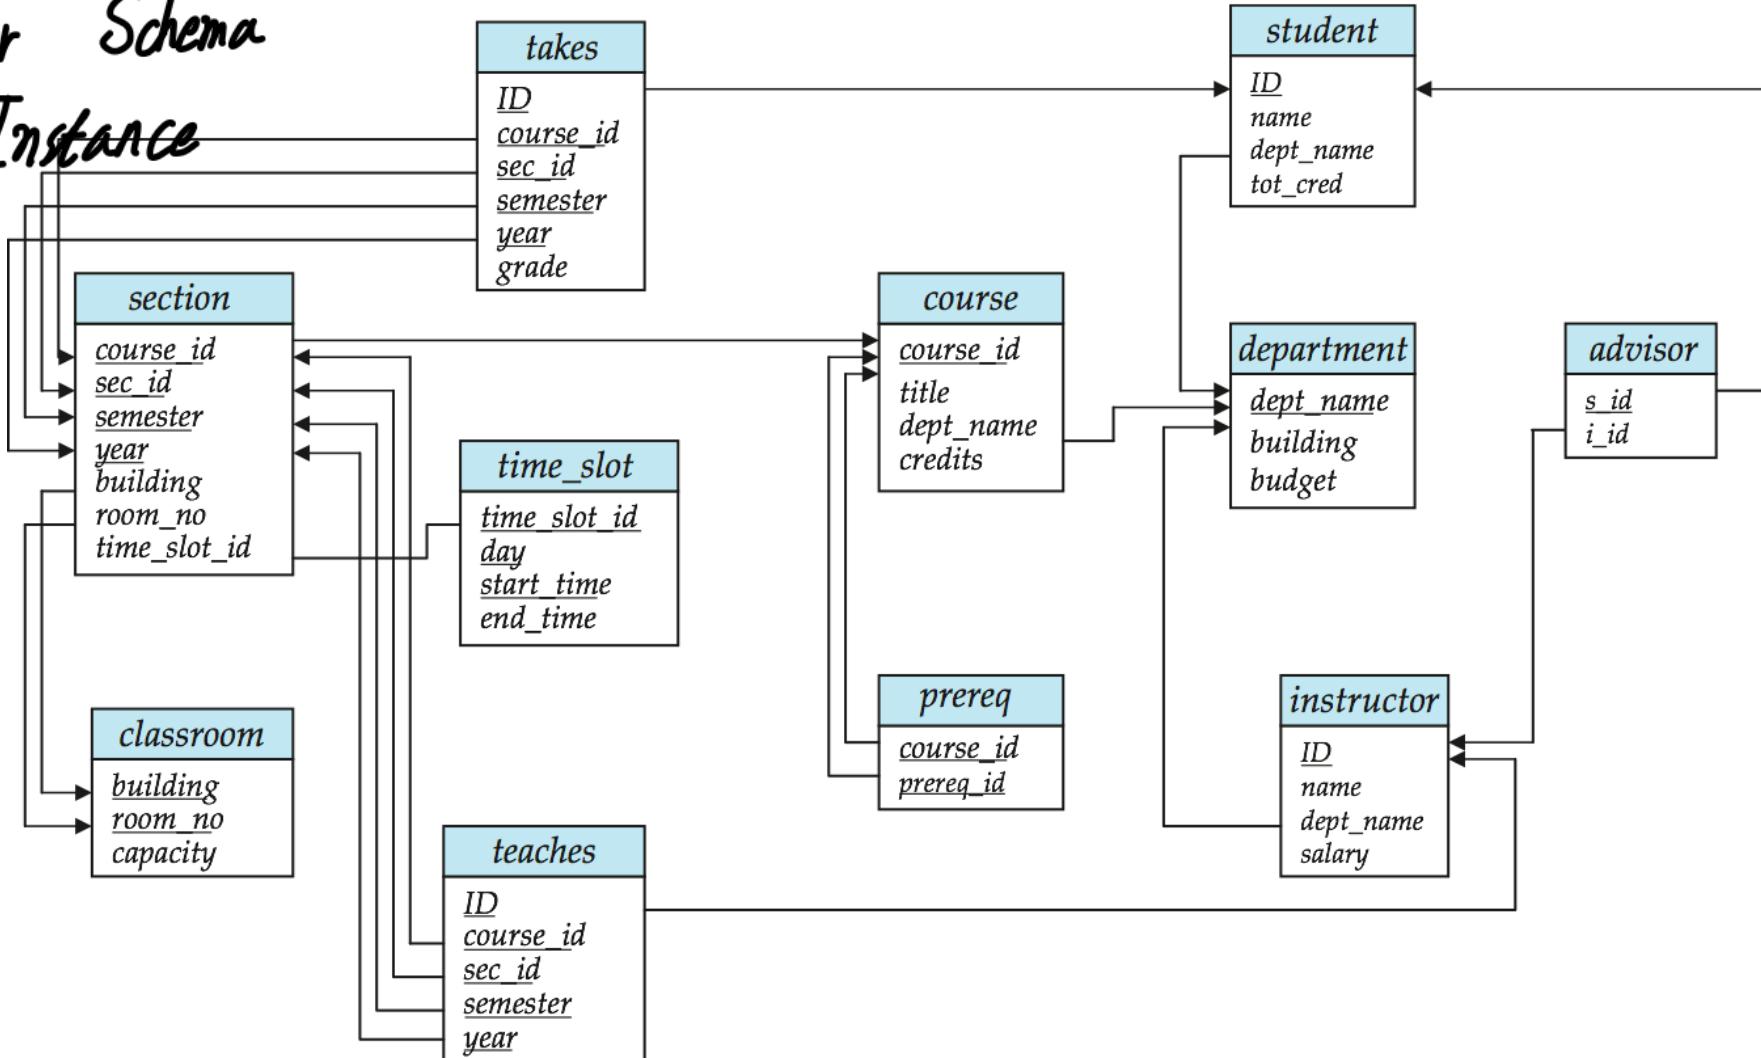
- $\sigma_{A=C}(r \times s)$

A	B	C	D	E
$\alpha$	1	$\alpha$	10	a
$\beta$	2	$\beta$	20	a
$\beta$	2	$\beta$	20	b

→ same

# Schema Diagram for University Database

Query for Schema  
not for Instance



# Example Queries

- Find all students in the CS department

$$\delta_{dept\_name = "CS"}(student)$$

- Find the name of each student in the CS department

$$\pi_{name}(\delta_{dept\_name = 'CS'}(student))$$

## Example Queries

- Find the names of all persons who are either an instructor or a student

$$\Pi_{name}(\text{student}) \cup \Pi_{name}(\text{instructor})$$

- Find the names of all persons who are both an instructor and a student  
(assuming names are unique)

$$\Pi_{name}(\text{student}) \cap \Pi_{name}(\text{instructor})$$

# Example Queries

- Find the names of all students who takes/took course CS-101.

Query 1

$$\Pi_{name} \left( \sigma_{course\_id='CS-101'} \left( \sigma_{takes.ID=Student.ID} (takes \times student) \right) \right)$$

$\hookrightarrow$  referring

Query 2

$$\Pi_{name} \left( \sigma_{takes.ID=student.ID} (student \times \sigma_{course.ID='CS-101'} (takes)) \right)$$

## Rename Operation In Cartesian Product

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.
- $\rho_N(E)$  returns the expression  $E$  under the name  $N$
- If a relational-algebra expression  $E$  has arity  $n$ , then

$$\rho_{N \underbrace{(A_1, A_2, \dots, A_n)}_{\text{new Attr Names}}}(E)$$

returns the result of expression  $E$  under the name  $N$ , and with the attributes renamed to  $A_1, A_2, \dots, A_n$ .

# Formal Definition

- A basic expression in the relational algebra consists of either one of the following:
  - A relation in the database
  - A constant relation
- Let  $E_1$  and  $E_2$  be relational-algebra expressions; the following are all *relational-algebra expressions*:
  - $E_1 \cup E_2$  Union
  - $E_1 - E_2$  difference
  - $E_1 \times E_2$  Cartesian product
  - $\sigma_P(E_1)$ ,  $P$  is a predicate on attributes in  $E_1$  select
  - $\Pi_S(E_1)$ ,  $S$  is a list consisting of some of the attributes in  $E_1$  project
  - $\rho_N(E_1)$ ,  $N$  is the new name for the result of  $E_1$  rename

$$\cap : \text{join?} \quad \bar{E}_1 - (\bar{E}_1 - \bar{E}_2) = E_1 \cap E_2$$

# Additional Operations

We define additional operations that do not add any power to the relational algebra, but that simplify common queries.

- Set intersection
- Natural join
- Assignment
- Outer Join
- Generalized Projection
- Aggregation

# Set-Intersection Operation

- Notation:  $r \cap s$

- Defined as:

$$r \cap s = \{ t \mid t \in r \text{ and } t \in s \}$$

- Assume union compatibility:

- $r, s$  have the *same arity*
  - attributes of  $r$  and  $s$  are compatible

- Note:  $r \cap s = r - (r - s)$

# Set-Intersection Operation – Example

- Relation  $r, s$ :

A	B
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$



A	B
$\alpha$	2
$\beta$	3

$s$

- $r \cap s$

A	B
$\alpha$	2

# Natural-Join Operation

- Let  $r(R)$  and  $s(S)$
- Notation:  $r \bowtie s$
- The result is a relation on schema  $\underline{R \cup S}$  which is obtained by considering each pair of tuples  $t_r$  from  $r$  and  $t_s$  from  $s$ .
  - If  $t_r$  and  $t_s$  have the same value on each of the attributes in  $\underline{R \cap S}$ , a tuple  $t$  is added to the result, where  $t$  has the same value as  $t_r$  on  $R$ , and  $t$  has the same value as  $t_s$  on  $S$ . ( i.e.,  $t[R]=t_r$  and  $t[S]=t_s$  )
- Example:

$$R = (A, B, C, D) \quad \& \quad S = (E, B, D)$$

Result schema =  $(A, B, C, D, E)$

$$r \bowtie s = \prod_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B=s.B \wedge r.D=s.D} (r \times s))$$

join table with primary-foreign key-set  
combine if match attr

# Natural-Join Operation – Example

- Relations  $r, s$ :

A	B	C	D
$\alpha$	1	$\alpha$	a
$\beta$	2	$\gamma$	a
$\gamma$	4	$\beta$	b
$\alpha$	1	$\gamma$	a
$\delta$	2	$\beta$	b

$r$

B	D	E
1	a	$\alpha$
3	a	$\beta$
1	a	$\gamma$
2	b	$\delta$
3	b	$\epsilon$

$s$

- $r \bowtie s$

A	B	C	D	E
$\alpha$	1	$\alpha$	a	$\alpha$
$\alpha$	1	$\alpha$	a	$\gamma$
$\alpha$	1	$\gamma$	a	$\alpha$
$\alpha$	1	$\gamma$	a	$\gamma$
$\delta$	2	$\beta$	b	$\delta$

# Properties

- $(r \bowtie s) \bowtie t = r \bowtie (s \bowtie t)$  : *exchange*

- If  $R \cap S = \emptyset$  then  $r \bowtie s = r \times s$

- If  $R = S$  then  $r \bowtie s = r \cap s$

- Theta Join

- combine selection with Cartesian product

$$r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$$

where  $\theta$  is a predicate on  $R \times S$

*select*

# Assignment Operation

- The assignment operation ( $\leftarrow$ )
  - provides a convenient way to express complex queries
  - write query as a sequential program consisting of a series of assignments
- Assignment must always be made to a temporary relation variable.
  - The result to the right hand side is assigned to the relation variable on the left hand side.
  - May use the variable in subsequent expressions.
- Example:  $r \cap s = r - (r - s)$

*temp  $\leftarrow r - s$       ) such as variable  
result  $\leftarrow r - temp$*

# Example Queries

- Find the names of all students who takes/took both courses CS-101 and CS-190.

- Query 1

$$\Pi_{name} (G_{course\_id='CS-101'}(\text{Student} \bowtie \text{takes}) \cap G_{course\_id='CS-190'}(\text{Student} \bowtie \text{takes}))$$

- Query 2

$$\Pi_{name} ((G_{course\_id='CS-101'}(\text{takes}) \cap G_{course\_id='CS-190'}(\text{takes})) \bowtie \text{student})$$

# Example Queries

- Find the IDs of all students who were taught by an instructor named Einstein in building 301.

```
temp1 ← σbuilding = '301' (Section) ∧ name = 'Einstein' (instructor) teaches
temp2 ← πcre_id, sec_id, smar, year(temp1)
result ← πID(temp2 ∧ takes)
Why? due to ID of instructor
```

```
temp1 ← section ∧ teaches
temp2 ← temp1 ∧ instructor
temp3 ← σname = 'Eisen', building = '301'(temp2)
temp4 ← πID(temp3 ∧ takes)
```

# **DISCUSSIONS – CHAPTER 2**

## Discussion 2-1

- A. Define *super key*, *candidate key*, and *primary key*.
- B. Why do you think “key”, without any qualifier, is meant to be candidate key instead of super key?

## Discussion 2-2

- Consider the foreign key constraint from the *dept\_name* attr. of *instructor* to the *department* relation. Give examples of inserts and deletes to these relations, which can cause a violation of the foreign key constraint.

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

*instructor*

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

*department*

## Discussion 2-3

- In the instance of *instructor* shown below, no two instructors have the same name. From this, can we conclude that *name* can be used as a super key (or primary key) of *instructor*?

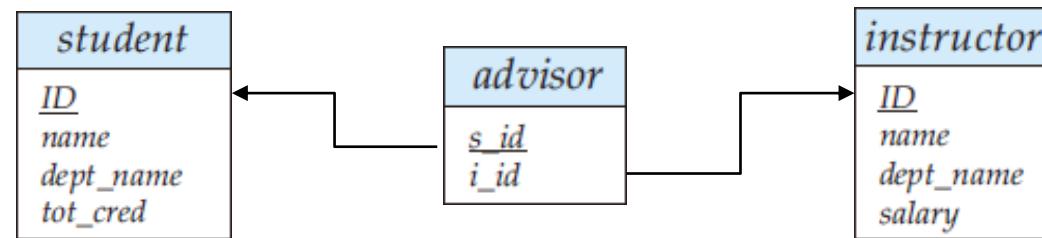
*Yes, but not primary key*

<i>instructor</i>	<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
	22222	Einstein	Physics	95000
	12121	Wu	Finance	90000
	32343	El Said	History	60000
	45565	Katz	Comp. Sci.	75000
	98345	Kim	Elec. Eng.	80000
	76766	Crick	Biology	72000
	10101	Srinivasan	Comp. Sci.	65000
	58583	Califieri	History	62000
	83821	Brandt	Comp. Sci.	92000
	15151	Mozart	Music	40000
	33456	Gold	Physics	87000
	76543	Singh	Finance	80000

## Discussion 2-4

- Consider the *advisor* relation shown below, with *s\_id* as the primary key of *advisor*. Suppose a student can have more than one advisor. Then, would *s\_id* still be a primary key of the *advisor* relation? If not, what should the primary key of *advisor* be?

*No, tuple of s\_id & i\_id*



## Discussion 2-5

*person* (pname, street, city)

*works* (pname, cname, salary)

*company* (cname, city) /\* keys are underlined

- Using the above database schema, represent the following queries in *relational algebra*.

A. Find all names of persons.

$\Pi_{pname}(\text{person})$

B. Find the names of persons who live in "Seoul"

$\Pi_{pname}(\sigma_{city = "Seoul"}(\text{person}))$

C. Find the names of persons who work in "SNU"

$\Pi_{pname}(\sigma_{cname = "SNU"}(\text{works}))$

## Discussion 2-6

*person* (pname, street, city)

*works* (pname, cname, salary)

*company* (cname, city) /\* keys are underlined

- Using the above database schema, represent the following queries in *relational algebra*.

A. Find all cities in the database

$\prod_{\text{city}}(\text{person}) \cup \prod_{\text{city}}(\text{company})$

B. Find the names of people who do not work

$\prod_{\text{pname}}(\text{person}) - \prod_{\text{pname}}(\text{works})$

C. Find names of people who work in "SNU" and earn more than 1,000,000

$\prod_{\text{pname}}(\sigma_{\text{cname} = "SNU" \wedge \text{salary} > 1,000,000}(\text{works}))$

## Discussion 2-7

*person* (pname, street, city)

*works* (pname, cname, salary)

*company* (cname, city) /\* keys are underlined

- Using the above database schema, represent the following queries in *relational algebra*.
  - Find names and addresses of persons who work for “SNU”.
  - Find company names located in the city where “SNU” is located.
  - Find names and addresses of persons who work for companies located in “Seoul”.

$$\begin{array}{c} \prod_{\text{street}, \text{pname}} \left( \text{person} \bowtie \left( \text{works} \bowtie \left( \text{city} = "Seoul" \bowtie \text{company} \right) \right) \right) \\ \prod_{\text{pname}, \text{cname}, \text{salary}} \end{array}$$

## Discussion 2-8

*person (pname, street, city)*

*works (pname, cname, salary)*

*company (cname, city)                    /\* keys are underlined*

- Using the above database schema, represent the following queries in *relational algebra*.
  - Find pairs of person names who live in the same city.*
  - Find pairs of person names who live in the same city, without duplicate pairs.*

## Discussion 2-9

Discuss why the following property holds for any two relations  $r(R)$  and  $s(S)$ .

- A.  $r \bowtie s = s \bowtie r$
- B. If  $R=S$  then  $r \bowtie s = r \cap s$ .

## Discussion 2-10

- Use parenthesis to indicate the proper orders of operations in the following relational algebra expression.

$$\sigma_p E_1 \cup E_2 \times E_3 \bowtie \rho_N E_3 \cap E_4 - E_5$$

## Discussion 2-11

- Fill the following precedence table for relational algebra operators:  $\sigma$ ,  $\cup$ ,  $\times$ ,  $\bowtie$ ,  $\rho$ ,  $\cap$ ,  $-$ ,  $\Pi$

Prec.	Operators	Notes
0	( )	highest
1		
2		
3		
4		

# Discussions 2-12

*person (pname, street, city)*

*works (pname, cname, salary)*

*company (cname, city)*

- Using the above database schema, represent the following queries in *relational algebra*.
  - *Find the name of person with the largest salary.*

**END OF CHAPTER 2**