

Intro to DB

CHAPTER 4

INTERMEDIATE SQL

Chapter 4: Intermediate SQL

- Join Expressions
- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

Joined Relations

- **Join operation**
 - takes two relations and return as a result another relation.
 - a Cartesian product which requires that tuples in the two relations match (under some condition).
 - also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

Join – On

- Join matches tuples with the same values specified in the on condition.

select *

from instructor join teaches on (instructor.ID=teaches.ID); \Rightarrow *from instructor, teaches
where instructor.ID = teaches.ID*

ID	name	dept_name	salary	ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
45565	Katz	Comp. Sci.	75000	22222	PHY-101	1	Fall	2009
76766	Crick	Biology	72000	32343	HIS-351	1	Spring	2010
76766	Crick	Biology	72000	45565	CS-101	1	Spring	2010
				45565	CS-319	1	Spring	2010
				76766	BIO-101	1	Summer	2009
				76766	BIO-301	1	Summer	2010

Join – On (cont.)

- Let $r_1(A_1, \dots, A_k), r_2(B_1, \dots, B_n)$
select *
from r_1 **join** r_2 **on** *condition*;

is equivalent to

select *
from r_1, r_2
where *condition*

- You can use expressions (that evaluate to relations) in **from** clauses
from E_1, \dots, E_n

select distinct *name, title*
from *instructor join teaches on (instructor.ID=teaches.ID)*, *course*
where *teaches.course_id=course.course_id*



Join Operations

- take two relations and return as a result another relation.
- typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join types	Join Conditions
inner join left outer join right outer join full outer join	natural on <u><predicate></u> using (A_1, A_1, \dots, A_n)

Inner join

- Inner join = join (default)

select * from instructor join teaches on (instructor.ID=teaches.ID);

= select * from instructor inner join teaches on (instructor.ID=teaches.ID);

↳ left only correspond result

- ON clause
 - an explicit join clause (acts like a **where** clause)
- USING clause
 - specifies which columns to test for *equality*
 - columns listed must be present in both of the two tables being joined

- Above example is equivalent to

select * from instructor inner join teaches using (ID);
A B column-list

Natural Join

- (Section 3.3.3)
- Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column

select *
from instructor **natural join** teaches;

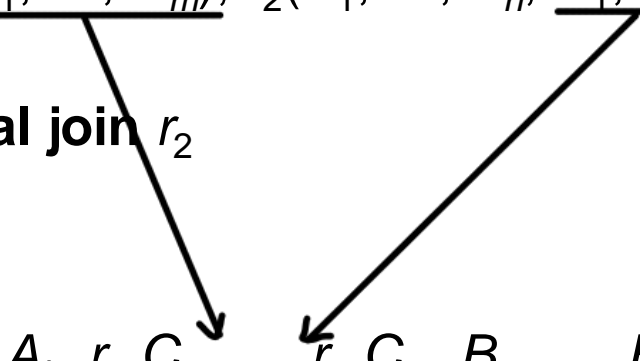
condition itself

- Let $r_1(A_1, \dots, A_k, \underline{C_1, \dots, C_m})$, $r_2(B_1, \dots, B_n, \underline{C_1, \dots, C_m})$

select *
from r_1 **natural join** r_2

is equivalent to

select $A_1, \dots, A_k, r_1.C_1, \dots, r_1.C_m, B_1, \dots, B_n$
from r_1, r_2
where $r_1.C_1=r_2.C_1$ **and** ... **and** $r_1.C_m=r_2.C_m$



Natural Join (cont.)

- Find the IDs of all students who were taught by an instructor named Einstein in building 301 (remove duplicates in the result).

```
select    distinct takes.ID
from      takes, section, teaches, instructor
where     takes.course_id=section.course_id and takes.sec_id=section.sec_id and
           takes.semester=section.semester and takes.year=section.year and
           section.course_id=teaches.course_id and section.sec_id=teaches.sec_id and
           section.year=teaches.year and section.year=teaches.year and
           teaches.ID=instructor.ID and
           instructor.name='Einstein' and building='301'
```

```
select    distinct takes.ID
from      (takes natural join section)
           join teaches using (crse_id, sec_id, smster, year)
           join instructor on (teaches.ID=instructor.ID) → due to takes.ID
where     instructor.name='Einstein' and building='301' (ambigious!)
```

Natural Join (cont.)

- **Danger in natural join:** beware of unrelated attributes with same name which get equated incorrectly

- List the names of instructors along with the titles of courses that they teach

- Incorrect version (makes course.dept_name = instructor.dept_name)

- **select distinct** name, title
from instructor natural join teaches **natural join** course;

it physic's dept instructor
teach electronic dept.

- Correct version

- **select distinct** name, title
from instructor **natural join** teaches, course
where teaches.course_id = course.course_id; (extra condition)

- Another correct version

- **select distinct** name, title
from (instructor **natural join** teaches)
join course **using** (course_id);

Join operation – Example

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

course **natural join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

Note: prereq information missing for CS-315 and
course information missing for CS-437.

(only left exist value)

Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

Left Outer Join : *preserve left table's dangling data*

course **natural left outer join** prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<u>null</u>

↳ unmatched, but survive

= course **left outer join** prereq **on** course.course_id = prereq.course_id

= course **left outer join** prereq **using** (course_id)

course

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Right Outer Join : *reverse of left*

course **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Full Outer Join : *left + right*

course **natural full outer join** prereq

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Joined Relations – Examples

course **inner join** *prereq* on

course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

course **left outer join** *prereq* on

course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

Summary – Joined Operations

- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using (A_1, A_1, \dots, A_n)

Integrity Constraints

- Integrity constraints guard against accidental damage to the database
 - by ensuring that authorized changes to the database do not result in a loss of data consistency
- Examples
 - A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number
- Constraints on a single relation
 - **not null**
 - **unique**
 - **primary key**
 - **check** (P), where P is a predicate
- Referential integrity

Data integrity : meet the rule of 'model'

↳ primary key - foreign key

Not Null Constraint

- Declare *name* for *budget* to be **not null**

name **varchar(20) not null**

budget **numeric(12,2) not null**

- Example

```
create table instructor (  
    ID          char(5) not null,  
    name        varchar(20) not null,  
    dept_name varchar(20),  
    salary     numeric(8,2))
```

The Unique & Primary Key Constraints

- **unique** (A_1, A_2, \dots, A_m)
 - A_1, A_2, \dots, A_m form a candidate key
 - can be null unless declared to be non null explicitly
- **primary key** (A_1, A_2, \dots, A_m)
 - A_1, A_2, \dots, A_m form the primary key for the table
 - Non null & unique

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null,  
    dept_name   varchar(20),  
    salary      numeric(8,2)  
    primary key (ID) );  
unique(ID)      ↪ primary key(ID, name)
```

The check clause


- check (P), where P is a predicate
- Example:
 - ensure that semester is one of fall, winter, spring or summer

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
)
```

*integrity constraints ↑
⇒ overhead also increase*

must be one of this lists

Referential Integrity

- “A value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.”
 - Example:
 - If “Biology” is a department name appearing in one of the tuples in the *course* relation,
 - then there exists a tuple in the *department* relation for “Biology”.
- specified as part of **create table** statement
- **Foreign Key**  $\Rightarrow == \text{check } (A \text{ in } B.\text{primary_key})$
 - the attributes that comprise the foreign key, and
 - the name of the relation referenced by the foreign key
 - By default, a foreign key references the primary key of the referenced table.


course_id	title	dept_name	credits
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

Cascading Actions in Referential Integrity

- **create table** *course* (
 course_id **char**(5) **primary key**,
 title **varchar**(20),
 dept_name **varchar**(20) **references** *department*
)
- automatically delete primary and link those key*
- **create table** *course* (
 ...
 dept_name **varchar**(20),
 foreign key (*dept_name*) **references** *department*
 on delete cascade
 ...
)
- automatic delete when primary key delete*
- alternative actions to cascade: **set null**, **set default**

Primary & Foreign Keys in Schema Diagram



Authorization Specification in SQL

- The **grant** statement is used to confer authorization
grant <privilege list>
on <relation name or view name>
to <user list>
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - a role (explained later)
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

Privileges in SQL

- **select**: allows read access to relation
 - (or the ability to query using the view)
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:

grant select on *branch* to U_1, U_2, U_3 *target users*

table

- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges
- **index**: allows creation and deletion of indices.
- **resources**: allows creation of new relations.
- **alteration**: allows addition or deletion of attributes in a relation.
- **drop**: allows deletion of relations.

Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.

revoke <privilege list>

on <relation name or view name>


from <user list>

- Example:

revoke select on *branch* from U_1, U_2, U_3

- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- All privileges that depend on the privilege being revoked are also revoked.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.

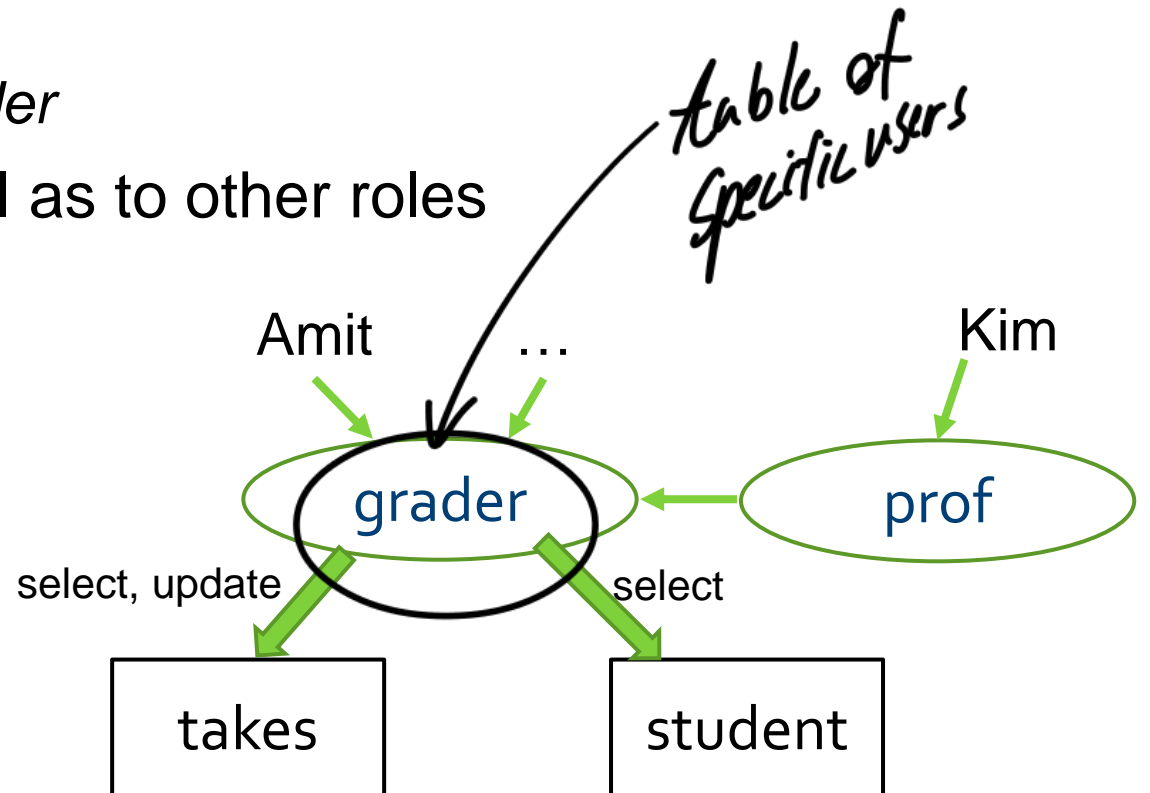
Authorization-Grant Graph



Roles

kind of authorized position

- Roles are used to represent a group of users and their privileges
 - **create role** *grader*;
- Privileges can be granted to roles:
 - **grant select on student to grader**
 - **grant select, update on takes to grader**
- Roles can be granted to users, as well as to other roles
 - **grant grader to Amit;**
 - **create role** *prof*;
 - **grant grader to prof;**
 - **grant prof to Kim;**




DISCUSSIONS – CHAPTER 4

Discussion 4-1

Write the following query in SQL using the **join** operator.

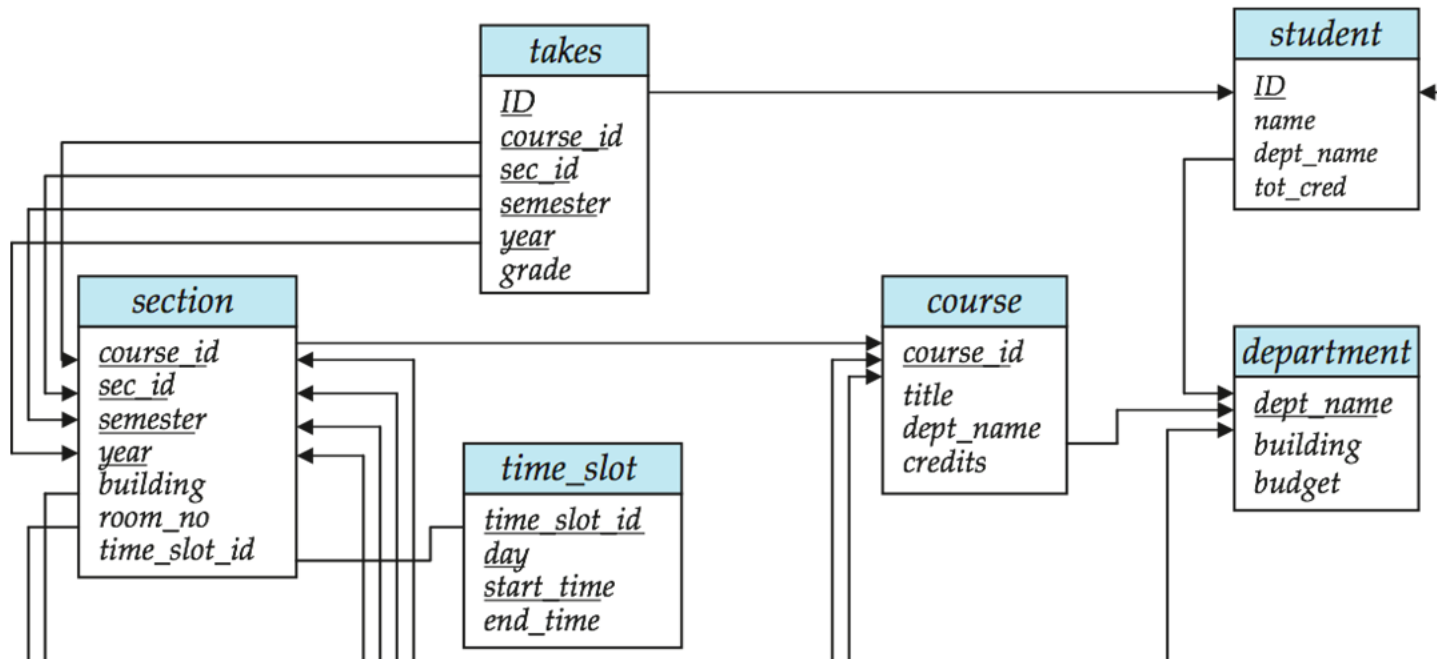
Find students (ID & name) who took a course in 2017.



Discussion 4-2

Write the following query in SQL using the **join** operator.


Find students (ID & name) who took a course in 2017 held in building '301'.



Discussion 4-3

Write the following query in SQL using the **join** operator.


Find students (ID & name) who took a course in 2017 offered by the 'CS' department.



Discussion 4-4

Write the following query in SQL.


List course titles with years they were offered. Include courses that were never offered.



Discussion 4-5

Write the following query in SQL.

Update the tot_cred of each student to the sum of credits that she/he has taken with a grade other than 'F' and null.



END OF CHAPTER 4